

Lab 7 - BCC406

REDES NEURAIS E APRENDIZAGEM EM PROFUNDIDADE

Modelos Generativos

Prof. Eduardo e Prof. Pedro

Objetivos:

- Parte I : Compressão com AE
- Parte II : Detecção de anomalias
- Parte III: Redes Generativas Adversariais

Data da entrega : 21/10

- Complete o código (marcado com **ToDo**) e quando requisitado, escreva textos diretamente nos notebooks. Onde tiver *None*, substitua pelo seu código.
- Execute todo notebook e salve tudo em um PDF **nomeado** como "NomeSobrenome-Lab.pdf"
- Envie o PDF via google [FORM](#)

Este notebook é baseado em tensorflow e Keras.

Parte I: Autoencoder para redução de dimensionalidade (30pt)

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import tensorflow as tf

from sklearn.metrics import accuracy_score, precision_s
from sklearn.model_selection import train_test_split
from tensorflow.keras import layers, losses
from tensorflow.keras.datasets import fashion_mnist
from tensorflow.keras.models import Model
```

Recursos X

Pretende obter mais memória e espaço em disco?

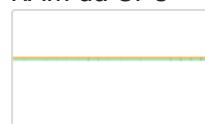
[Atualizar para o Colab Pro](#)

(GPU) de back-end do Google Compute Engine em Python 3
A mostrar os recursos desde 14:10...

RAM do sistema



RAM da GPU



Disco



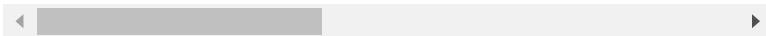
Carrega dataset Fashion MNIST dataset. Cada imagem tem resolução 28x28 pixels.

```
(x_train, _), (x_test, _) = fashion_mnist.load_data()

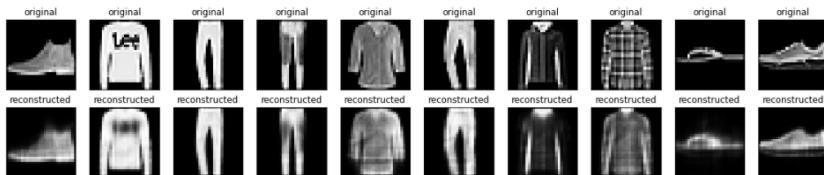
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.

print (x_train.shape)
print (x_test.shape)

Downloading data from https://storage.googleapis.com
32768/29515 [=====] -
40960/29515 [=====]
Downloading data from https://storage.googleapis.com
26427392/26421880 [=====
26435584/26421880 [=====
Downloading data from https://storage.googleapis.com
16384/5148 [=====
Downloading data from https://storage.googleapis.com
4423680/4422102 [=====] .
4431872/4422102 [=====] .
(60000, 28, 28)
(10000, 28, 28)
```



▼ Exemplo de classes



Abaixo exemplo de implementação de autoencoder apenas com camadas densas. O encoder, comprime as imagens em 4 dimensões (latent_dim), e o decoder reconstrói a imagem a partir do vetor latente.

O exemplo abaixo usa a [Keras Model Subclassing API](#).

```
latent_dim = 4

class Autoencoder(Model):
    def __init__(self, latent_dim):
        super(Autoencoder, self).__init__()
        self.latent_dim = latent_dim
        self.encoder = tf.keras.Sequential([
            layers.Flatten(),
            layers.Dense(latent_dim, activation='relu'),
```

```
])
self.decoder = tf.keras.Sequential([
    layers.Dense(784, activation='sigmoid'),
    layers.Reshape((28, 28))
])

def call(self, x):
    encoded = self.encoder(x)
    decoded = self.decoder(encoded)
    return decoded

autoencoder = Autoencoder(latent_dim)

autoencoder.compile(optimizer='adam', loss=losses.MeanS

autoencoder.fit(x_train, x_train,
                 epochs=10,
                 shuffle=True,
                 validation_data=(x_test, x_test))

Epoch 1/10
1875/1875 [=====] - 8s 2r
Epoch 2/10
1875/1875 [=====] - 4s 2r
Epoch 3/10
1875/1875 [=====] - 4s 2r
Epoch 4/10
1875/1875 [=====] - 4s 2r
Epoch 5/10
1875/1875 [=====] - 4s 2r
Epoch 6/10
1875/1875 [=====] - 4s 2r
Epoch 7/10
1875/1875 [=====] - 4s 2r
Epoch 8/10
1875/1875 [=====] - 4s 2r
Epoch 9/10
1875/1875 [=====] - 4s 2r
Epoch 10/10
1875/1875 [=====] - 4s 2r
<keras.callbacks.History at 0x7f39a0336d90>
```



Treine o modelo e veja os resultados da re-construção.

```
encoded_imgs = autoencoder.encoder(x_test).numpy()
decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()

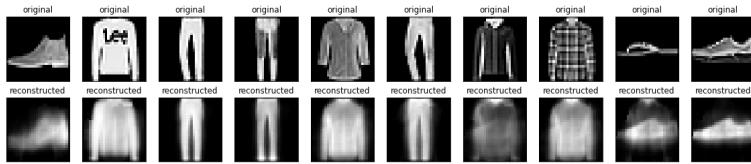
n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
    # display original
```

```

ax = plt.subplot(2, n, i + 1)
plt.imshow(x_test[i])
plt.title("original")
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)

# display reconstruction
ax = plt.subplot(2, n, i + 1 + n)
plt.imshow(decoded_imgs[i])
plt.title("reconstructed")
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()

```



▼ ToDo : Testes (15pt)

Faça testes com vetor latente de dimensões 2, 8, 16 e 64.

2 Dimensões

```

latent_dim = 2 # Definir tamanho do vetor Latente

# Criar e treinar o Auto-Encoder
autoencoder = Autoencoder(latent_dim)
autoencoder.compile(optimizer='adam', loss=losses.MeanS
autoencoder.fit(x_train, x_train,
                 epochs=10,
                 shuffle=True,
                 validation_data=(x_test, x_test))

# Realizar Compressões e Descompressões nas imagens da
encoded_imgs = autoencoder.encoder(x_test).numpy()
decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()

# Plotar 10 primeiros resultados
n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
    # display original

```

```

ax = plt.subplot(2, n, i + 1)
plt.imshow(x_test[i])
plt.title("original")
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)

# display reconstruction
ax = plt.subplot(2, n, i + 1 + n)
plt.imshow(decoded_imgs[i])
plt.title("reconstructed")
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()

Epoch 1/10
1875/1875 [=====] - 4s 2r
Epoch 2/10
1875/1875 [=====] - 4s 2r
Epoch 3/10
1875/1875 [=====] - 4s 2r
Epoch 4/10
1875/1875 [=====] - 4s 2r
Epoch 5/10
1875/1875 [=====] - 4s 2r
Epoch 6/10
1875/1875 [=====] - 4s 2r
Epoch 7/10
1875/1875 [=====] - 4s 2r
Epoch 8/10
1875/1875 [=====] - 4s 2r
Epoch 9/10
1875/1875 [=====] - 4s 2r
Epoch 10/10
1875/1875 [=====] - 4s 2r

```

| original |
|---|---|---|---|---|---|---|---|---|--|
|  |  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |  |



8 Dimensões

```

latent_dim = 8 # Definir tamanho do vetor Latente

# Criar e treinar o Auto-Encoder
autoencoder = Autoencoder(latent_dim)
autoencoder.compile(optimizer='adam', loss=losses.MeanS
autoencoder.fit(x_train, x_train,
                 epochs=10,

```

```
shuffle=True,  
validation_data=(x_test, x_test))  
  
# Realizar Compressões e Descompressões nas imagens da  
encoded_imgs = autoencoder.encoder(x_test).numpy()  
decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()  
  
# Plotar 10 primeiros resultados  
n = 10  
plt.figure(figsize=(20, 4))  
for i in range(n):  
    # display original  
    ax = plt.subplot(2, n, i + 1)  
    plt.imshow(x_test[i])  
    plt.title("original")  
    plt.gray()  
    ax.get_xaxis().set_visible(False)  
    ax.get_yaxis().set_visible(False)  
  
    # display reconstruction  
    ax = plt.subplot(2, n, i + 1 + n)  
    plt.imshow(decoded_imgs[i])  
    plt.title("reconstructed")  
    plt.gray()  
    ax.get_xaxis().set_visible(False)  
    ax.get_yaxis().set_visible(False)  
plt.show()
```

Epoch 1/10

16 Dimensões

```
latent_dim = 16 # Definir tamanho do vetor Latente

# Criar e treinar o Auto-Encoder
autoencoder = Autoencoder(latent_dim)
autoencoder.compile(optimizer='adam', loss=losses.MeanS
autoencoder.fit(x_train, x_train,
                 epochs=10,
                 shuffle=True,
                 validation_data=(x_test, x_test))

# Realizar Compressões e Descompressões nas imagens da
encoded_imgs = autoencoder.encoder(x_test).numpy()
decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()

# Plotar 10 primeiros resultados
n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
    # display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i])
    plt.title("original")
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # display reconstruction
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i])
    plt.title("reconstructed")
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```

```
Epoch 1/10
1875/1875 [=====] - 4s 2r
Epoch 2/10
1875/1875 [=====] - 4s 2r
Epoch 3/10
1875/1875 [=====] - 4s 2r
Epoch 4/10
1875/1875 [=====] - 4s 2r
Epoch 5/10
1875/1875 [=====] - 4s 2r
Epoch 6/10
1875/1875 [=====] - 4s 2r
Epoch 7/10
1875/1875 [=====] - 4s 2r
```

64 Dimensões

```
Epoch 8/10
```

```
latent_dim = 64 # Definir tamanho do vetor Latente

# Criar e treinar o Auto-Encoder
autoencoder = Autoencoder(latent_dim)
autoencoder.compile(optimizer='adam', loss=losses.MeanS
autoencoder.fit(x_train, x_train,
                 epochs=10,
                 shuffle=True,
                 validation_data=(x_test, x_test))

# Realizar Compressões e Descompressões nas imagens da
encoded_imgs = autoencoder.encoder(x_test).numpy()
decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()

# Plotar 10 primeiros resultados
n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
    # display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i])
    plt.title("original")
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # display reconstruction
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i])
    plt.title("reconstructed")
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```

```

Epoch 1/10
1875/1875 [=====] - 4s 2r
Epoch 2/10
1875/1875 [=====] - 4s 2r
Epoch 3/10
1875/1875 [=====] - 4s 2r
Epoch 4/10
1875/1875 [=====] - 4s 2r
Epoch 5/10
1875/1875 [=====] - 4s 2r
Epoch 6/10
1875/1875 [=====] - 4s 2r
Epoch 7/10
1875/1875 [=====] - 4s 2r
Epoch 8/10
1875/1875 [=====] - 4s 2r
Epoch 9/10
1875/1875 [=====] - 4s 2r
Epoch 10/10
1875/1875 [=====] - 4s 2r

```

original	reconstructed								

▼ ToDo : Responda (15pt)

Escreva suas conclusões sobre os testes executados

Quando é aumentado o número de dimensões do vetor latente, a imagem decodificada fica mais parecida com a imagem original.

Parte II: Detecção de anomalias

▼ (30pt)

Intro

Neste exemplo, você vai detectar anomalias em sinais de eletrocardiograma (ECG). Para tal, treine um autoencoder no dataset [ECG5000 dataset](#). Este dataset contém 5000 batimentos de ECG (<https://en.wikipedia.org/wiki/Electrocardiography>), cada um com 140 amostras (pontos) na curva. Cada instância

da base de dados (um batimento) foi rotulado como zero (0) ou um (1). A classe zero corresponde a um batimento anormal e a classe um a um batimento de classe normal. Queremos identificar os anormais.

Para detectar anomalias usando um autoencoder você deve treinar um autoencoder apenas em batimentos normais. Ele vai aprender a re-construir os batimentos saudáveis. A hipótese é que os batimentos anormais vão divergir do padrão, quando compararmos a entrada com a re-construção.

▼ Carrega base de ECG

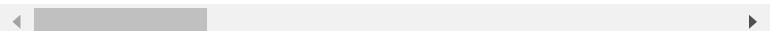
Base de dados detalhada no site:

timeseriesclassification.com.

```
# Download the dataset
dataframe = pd.read_csv('http://storage.googleapis.com/
raw_data = dataframe.values
dataframe.head()
```

	0	1	2	3	
0	-0.112522	-2.827204	-3.773897	-4.349751	-4.376
1	-1.100878	-3.996840	-4.285843	-4.506579	-4.022
2	-0.567088	-2.593450	-3.874230	-4.584095	-4.187
3	0.490473	-1.914407	-3.616364	-4.318823	-4.268
4	0.800232	-0.874252	-2.384761	-3.973292	-4.338

5 rows × 141 columns



```
# The last element contains the labels
labels = raw_data[:, -1]

# The other data points are the electrocardiogram data
data = raw_data[:, 0:-1]

train_data, test_data, train_labels, test_labels = train_
data, labels, test_size=0.2, random_state=21
)
```

Normaliza entre $[0,1]$.

```
min_val = tf.reduce_min(train_data)
max_val = tf.reduce_max(train_data)

train_data = (train_data - min_val) / (max_val - min_val)
test_data = (test_data - min_val) / (max_val - min_val)

train_data = tf.cast(train_data, tf.float32)
test_data = tf.cast(test_data, tf.float32)
```

Vamos separar os batimentos normais (label 1) para treinar o Autoencoder.

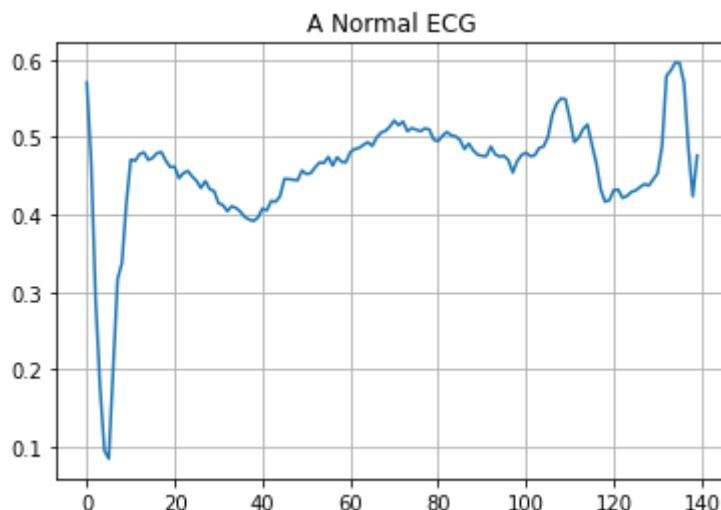
```
train_labels = train_labels.astype(bool)
test_labels = test_labels.astype(bool)

normal_train_data = train_data[train_labels]
normal_test_data = test_data[test_labels]

anomalous_train_data = train_data[~train_labels]
anomalous_test_data = test_data[~test_labels]
```

Plote um batimento normal.

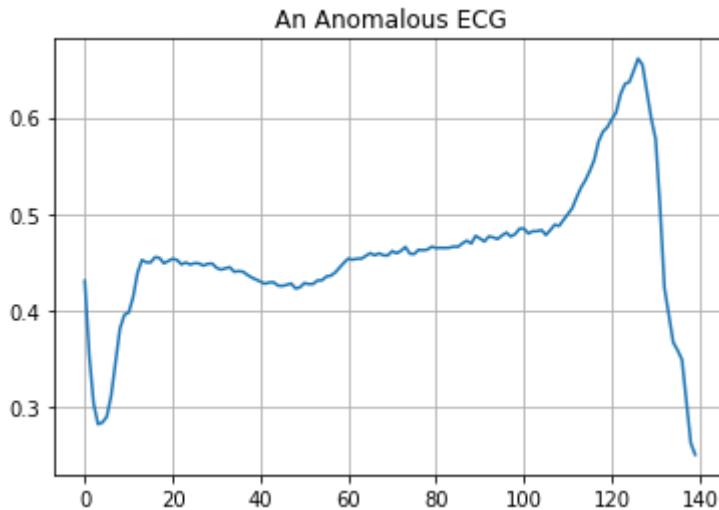
```
plt.grid()
plt.plot(np.arange(140), normal_train_data[0])
plt.title("A Normal ECG")
plt.show()
```



Plote um batimento anômalo.

```
plt.grid()
```

```
plt.plot(np.arange(140), anomalous_train_data[0])
plt.title("An Anomalous ECG")
plt.show()
```



▼ ToDo : Construção de um modelo (30pt)

Construa um modelo. Primeiramente tente construir apenas com camadas densas. Depois, tente construir um modelo com camadas de convolução de uma dimensão (Lembre-se que um sinal de ECG é uma série temporal de uma dimensão). [Conv1D](#)

```
# Importações
from tensorflow.keras.layers import Conv1D
from tensorflow.keras.layers import Input

class AnomalyDetector(Model):
    def __init__(self):
        super(AnomalyDetector, self).__init__()
        #self.encoder = tf.keras.Sequential([
        #    # Todo: crie o encoder, com um gargalo
        self.encoder = tf.keras.Sequential([
            Conv1D(filters=
            Conv1D(filters=
            layers.Flatten(
            layers.Dense(52
            layers.Dense(1,
            layers.Dense(52

#self.decoder = tf.keras.Sequential([
    # Todo: crie as camadas do decoder
    self.decoder = tf.keras.Sequential([layers.Flatten(
        layers.Dense(52
        layers.Dense(14

    def call(self, x):
```

```
encoded = self.encoder(x)
decoded = self.decoder(encoded)
return decoded

autoencoder2 = AnomalyDetector()

class AnomalyDetector(Model):
    def __init__(self):
        super(AnomalyDetector, self).__init__()
        #self.encoder = tf.keras.Sequential([
        #    # Todo: crie o encoder, com um gargalo
        self.encoder = tf.keras.Sequential([layers.Flatten(
            layers.Dense(52,
            layers.Dense(1,
            layers.Dense(52

#self.decoder = tf.keras.Sequential([
#    # Todo: crie as camadas do decoder
    self.decoder = tf.keras.Sequential([layers.Dense(14
def call(self, x):
    encoded = self.encoder(x)
    decoded = self.decoder(encoded)
    return decoded

autoencoder = AnomalyDetector()

autoencoder.compile(optimizer='adam', loss='mae')

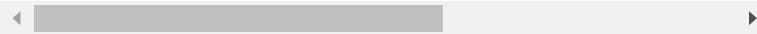
autoencoder2.compile(optimizer='adam', loss='mae')
```

Depois de treinar com os batimentos normais, avalie com os anormais.

```
history = autoencoder.fit(normal_train_data, normal_train_data,
                           epochs=20,
                           batch_size=512,
                           validation_data=(test_data, test_data),
                           shuffle=True)

Epoch 1/20
5/5 [=====] - 1s 36ms/step
Epoch 2/20
5/5 [=====] - 0s 8ms/step
Epoch 3/20
5/5 [=====] - 0s 8ms/step
Epoch 4/20
5/5 [=====] - 0s 9ms/step
Epoch 5/20
5/5 [=====] - 0s 9ms/step
Epoch 6/20
5/5 [=====] - 0s 8ms/step
Epoch 7/20
```

```
5/5 [=====] - 0s 8ms/step
Epoch 8/20
5/5 [=====] - 0s 9ms/step
Epoch 9/20
5/5 [=====] - 0s 8ms/step
Epoch 10/20
5/5 [=====] - 0s 8ms/step
Epoch 11/20
5/5 [=====] - 0s 9ms/step
Epoch 12/20
5/5 [=====] - 0s 8ms/step
Epoch 13/20
5/5 [=====] - 0s 12ms/step
Epoch 14/20
5/5 [=====] - 0s 8ms/step
Epoch 15/20
5/5 [=====] - 0s 8ms/step
Epoch 16/20
5/5 [=====] - 0s 8ms/step
Epoch 17/20
5/5 [=====] - 0s 8ms/step
Epoch 18/20
5/5 [=====] - 0s 8ms/step
Epoch 19/20
5/5 [=====] - 0s 8ms/step
Epoch 20/20
5/5 [=====] - 0s 11ms/step
```



```
history2 = autoencoder2.fit(np.expand_dims(normal_train,
                                             epochs=20,
                                             batch_size=512,
                                             validation_data=(np.expand_dims(test_data, axis=-1),
                                                               shuffle=True))
```

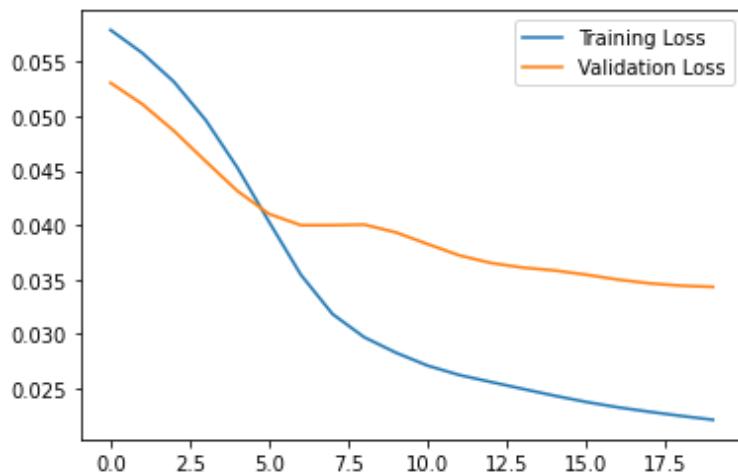
```
Epoch 1/20
5/5 [=====] - 9s 89ms/step
Epoch 2/20
5/5 [=====] - 0s 22ms/step
Epoch 3/20
5/5 [=====] - 0s 21ms/step
Epoch 4/20
5/5 [=====] - 0s 21ms/step
Epoch 5/20
5/5 [=====] - 0s 18ms/step
Epoch 6/20
5/5 [=====] - 0s 18ms/step
Epoch 7/20
5/5 [=====] - 0s 18ms/step
Epoch 8/20
5/5 [=====] - 0s 18ms/step
Epoch 9/20
5/5 [=====] - 0s 20ms/step
Epoch 10/20
5/5 [=====] - 0s 17ms/step
Epoch 11/20
5/5 [=====] - 0s 19ms/step
```

```
Epoch 12/20
5/5 [=====] - 0s 17ms/step
Epoch 13/20
5/5 [=====] - 0s 17ms/step
Epoch 14/20
5/5 [=====] - 0s 17ms/step
Epoch 15/20
5/5 [=====] - 0s 17ms/step
Epoch 16/20
5/5 [=====] - 0s 17ms/step
Epoch 17/20
5/5 [=====] - 0s 19ms/step
Epoch 18/20
5/5 [=====] - 0s 18ms/step
Epoch 19/20
5/5 [=====] - 0s 18ms/step
Epoch 20/20
5/5 [=====] - 0s 17ms/step
```

◀ ▶

```
plt.plot(history.history["loss"], label="Training Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
plt.legend()
```

<matplotlib.legend.Legend at 0x7f392062a690>



```
plt.plot(history2.history["loss"], label="Training Loss")
plt.plot(history2.history["val_loss"], label="Validation Loss")
plt.legend()
```

```
<matplotlib.legend.Legend at 0x7f3920502dd0>
```

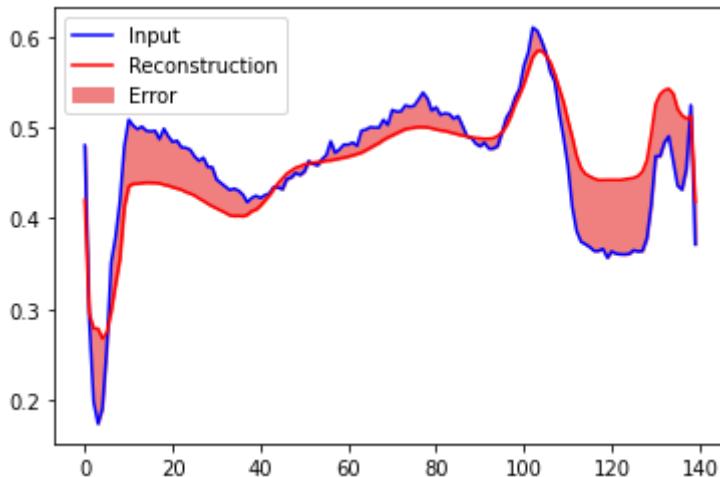


Você vai considerar um batimento como anômalo se ele divergir mais que um desvio padrão das amostras normais. Primeiro, vamos plotar um batimento normal a partir da base de treino e sua reconstrução. Assim, poderemos calcular o erro de re-construção.



```
encoded_data = autoencoder.encoder(normal_test_data).numpy()
decoded_data = autoencoder.decoder(encoded_data).numpy()

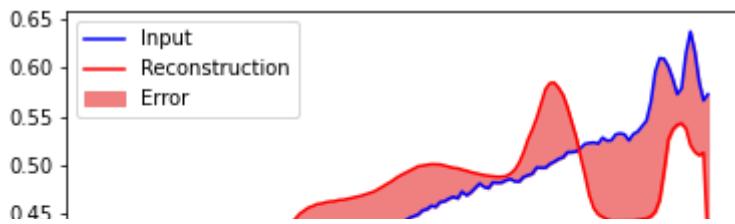
plt.plot(normal_test_data[0], 'b')
plt.plot(decoded_data[0], 'r')
plt.fill_between(np.arange(140), decoded_data[0], normal_test_data[0], color='red', alpha=0.2)
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()
```



Vamos fazer o mesmo para um batimento anômalo.

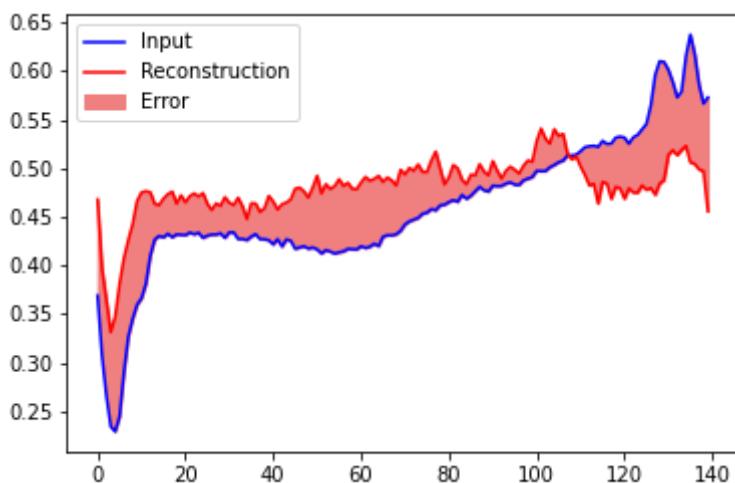
```
encoded_data = autoencoder.encoder(anomalous_test_data)
decoded_data = autoencoder.decoder(encoded_data).numpy()

plt.plot(anomalous_test_data[0], 'b')
plt.plot(decoded_data[0], 'r')
plt.fill_between(np.arange(140), decoded_data[0], anomalous_test_data[0], color='red', alpha=0.2)
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()
```



```
encoded_data2 = autoencoder2.encoder(np.expand_dims(ano
decoded_data2 = autoencoder2.decoder(encoded_data).nump
```

```
plt.plot(anomalous_test_data[0], 'b')
plt.plot(decoded_data2[0], 'r')
plt.fill_between(np.arange(140), decoded_data2[0], anom
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()
```



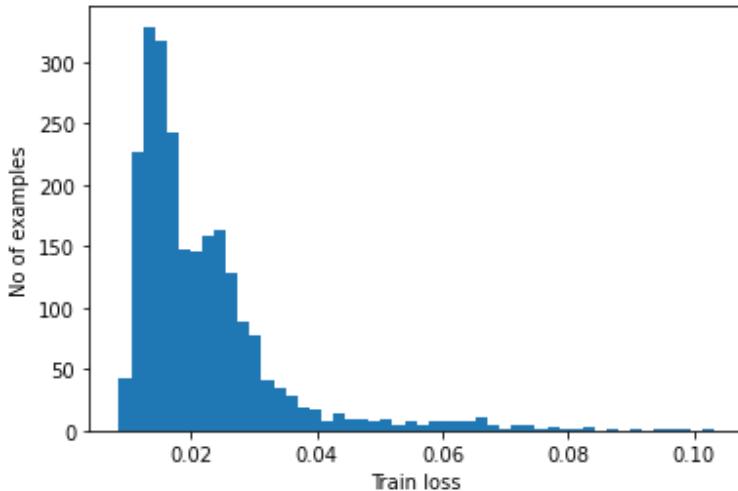
▼ Detectando as anomalias

Vamos detectar as anomalias se o erro de reconstrução for maior que um limiar. Aqui, vamos calcular o erro médio para os exemplos normais do treino e depois, classificar os anormais do teste, que tenha erro de reconstrução maior que um desvio padrão.

Plota erro de reconstrução de batimentos normais do treino

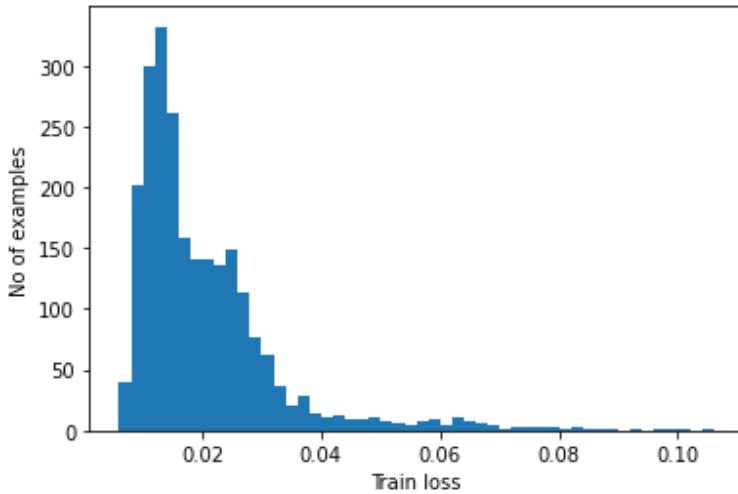
```
reconstructions = autoencoder.predict(normal_train_data
train_loss = tf.keras.losses.mae(reconstructions, norma

plt.hist(train_loss[None,:], bins=50)
plt.xlabel("Train loss")
plt.ylabel("No of examples")
plt.show()
```



```
reconstructions2 = autoencoder2.predict(np.expand_dims(  
train_loss2 = tf.keras.losses.mae(reconstructions2, nor
```

```
plt.hist(train_loss2[None,:], bins=50)  
plt.xlabel("Train loss")  
plt.ylabel("No of examples")  
plt.show()
```



Escolha do limiar.

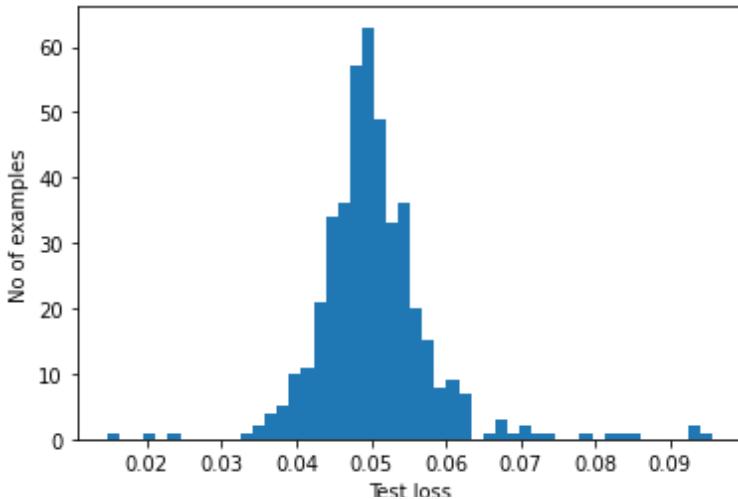
```
threshold = np.mean(train_loss) + np.std(train_loss)  
print("Threshold: ", threshold)
```

```
Threshold:  0.034146238
```

```
reconstructions = autoencoder.predict(anomalous_test_da  
test_loss = tf.keras.losses.mae(reconstructions, anomal
```

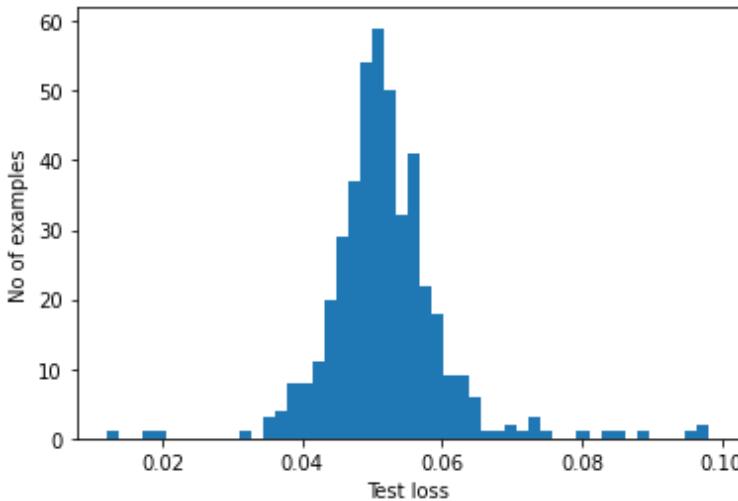
```
plt.hist(test_loss[None, :], bins=50)  
plt.xlabel("Test loss")
```

```
plt.ylabel("No of examples")
plt.show()
```



```
reconstructions2 = autoencoder2.predict(np.expand_dims(
test_loss2 = tf.keras.losses.mae(reconstructions2, anom
```

```
plt.hist(test_loss2[None, :], bins=50)
plt.xlabel("Test loss")
plt.ylabel("No of examples")
plt.show()
```



Classificação.

```
def predict(model, data, threshold):
    reconstructions = model(data)
    loss = tf.keras.losses.mae(reconstructions, data)
    return tf.math.less(loss, threshold)

def print_stats(predictions, labels):
    print("Accuracy = {}".format(accuracy_score(labels, predictions)))
    print("Precision = {}".format(precision_score(labels, predictions)))
    print("Recall = {}".format(recall_score(labels, predictions)))
```

Calcule a acurácia para os dois modelos (com camadas densas e convolucionais)

```
preds = predict(autoencoder, test_data, threshold)
print_stats(preds, test_labels)

Accuracy = 0.945
Precision = 0.9922027290448343
Recall = 0.9089285714285714

preds = predict(autoencoder2, np.expand_dims(test_data,

predito = []
for i in range(preds.shape[0]):
    predito.append(bool(preds[i][0]))

print_stats(predito, test_labels)

Accuracy = 0.948
Precision = 0.9922480620155039
Recall = 0.9142857142857143
```

Parte III: Redes Generativas

▼ Adversariais (40pt)

Leia o tutorial sobre a pix2pix em [Tensorflow Tutorials](#). O pix2pix foi apresentado em [Image-to-image translation with conditional adversarial networks by Isola et al. \(2017\)](#) e se trata de uma rede generativa adversarial condicional para geração de fachadas de prédios condicionada a uma máscara representando a arquitetura. baixe o notebook do tutorial, estude e treine a GAN. Após o treinamento, construa você mesmo 3 máscaras (usando algum software de desenho) e faça uma inferência com a rede. Anexe no notebook a máscara e sua respectiva saída.

▼ ToDo : Fachadas de prédios (40pt)

Importações

```
import tensorflow as tf
```

```
import os
import pathlib
import time
import datetime

from matplotlib import pyplot as plt
from IPython import display
```

Carregar o conjunto de dados

```
dataset_name = "facades"

_URL = f'http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/facades/facades.tar.gz'
```

```
path_to_zip = tf.keras.utils.get_file(
    fname=f'{dataset_name}.tar.gz',
    origin=_URL,
    extract=True)
```

```
path_to_zip = pathlib.Path(path_to_zip)
```

```
PATH = path_to_zip.parent/dataset_name
```

```
Downloading data from http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/facades/facades.tar.gz
30171136/30168306 [=====]
30179328/30168306 [=====]
```



```
list(PATH.parent.iterdir())
```

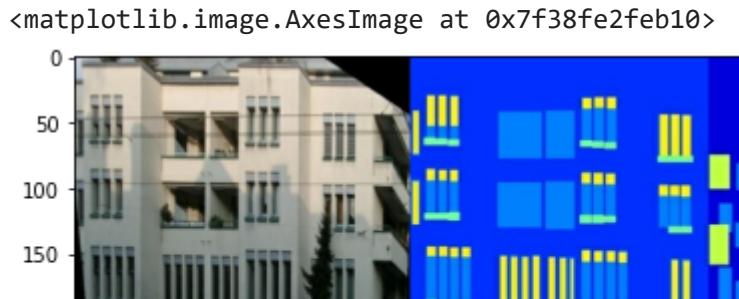
```
[PosixPath('/root/.keras/datasets/fashion-mnist'),
 PosixPath('/root/.keras/datasets/facades.tar.gz'),
 PosixPath('/root/.keras/datasets/facades')]
```



```
sample_image = tf.io.read_file(str(PATH / 'train/1.jpg'))
sample_image = tf.io.decode_jpeg(sample_image)
print(sample_image.shape)
```

```
(256, 512, 3)
```

```
plt.figure()
plt.imshow(sample_image)
```



Separar imagens reais das fachadas



```
def load(image_file):
    # Read and decode an image file to a uint8 tensor
    image = tf.io.read_file(image_file)
    image = tf.io.decode_jpeg(image)

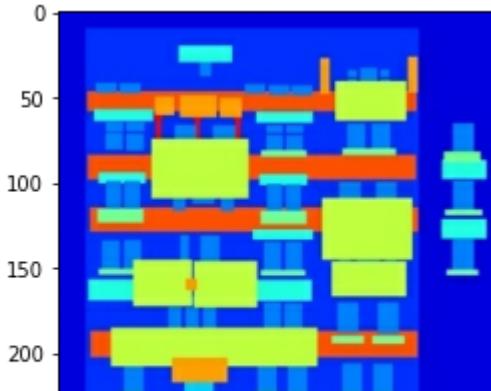
    # Split each image tensor into two tensors:
    # - one with a real building facade image
    # - one with an architecture label image
    w = tf.shape(image)[1]
    w = w // 2
    input_image = image[:, w:, :]
    real_image = image[:, :w, :]

    # Convert both images to float32 tensors
    input_image = tf.cast(input_image, tf.float32)
    real_image = tf.cast(real_image, tf.float32)

    return input_image, real_image

inp, re = load(str(PATH / 'train/100.jpg'))
# Casting to int for matplotlib to display the images
plt.figure()
plt.imshow(inp / 255.0)
plt.figure()
plt.imshow(re / 255.0)
```

```
<matplotlib.image.AxesImage at 0x7f38b6ba3450>
```



```
# The facade training set consist of 400 images
BUFFER_SIZE = 400
# The batch size of 1 produced better results for the U
BATCH_SIZE = 1
# Each image is 256x256 in size
IMG_WIDTH = 256
IMG_HEIGHT = 256
```



```
def resize(input_image, real_image, height, width):
    input_image = tf.image.resize(input_image, [height, width],
                                  method=tf.image.ResizeMethod.NEAREST_NEIGHBOR)
    real_image = tf.image.resize(real_image, [height, width],
                                method=tf.image.ResizeMethod.NEAREST_NEIGHBOR)

    return input_image, real_image
```

```
def random_crop(input_image, real_image):
    stacked_image = tf.stack([input_image, real_image], axis=0)
    cropped_image = tf.image.random_crop(
        stacked_image, size=[2, IMG_HEIGHT, IMG_WIDTH, 3])

    return cropped_image[0], cropped_image[1]
```

```
# Normalizing the images to [-1, 1]
def normalize(input_image, real_image):
    input_image = (input_image / 127.5) - 1
    real_image = (real_image / 127.5) - 1

    return input_image, real_image
```

```
@tf.function()
def random_jitter(input_image, real_image):
    # Resizing to 286x286
    input_image, real_image = resize(input_image, real_image, 286, 286)

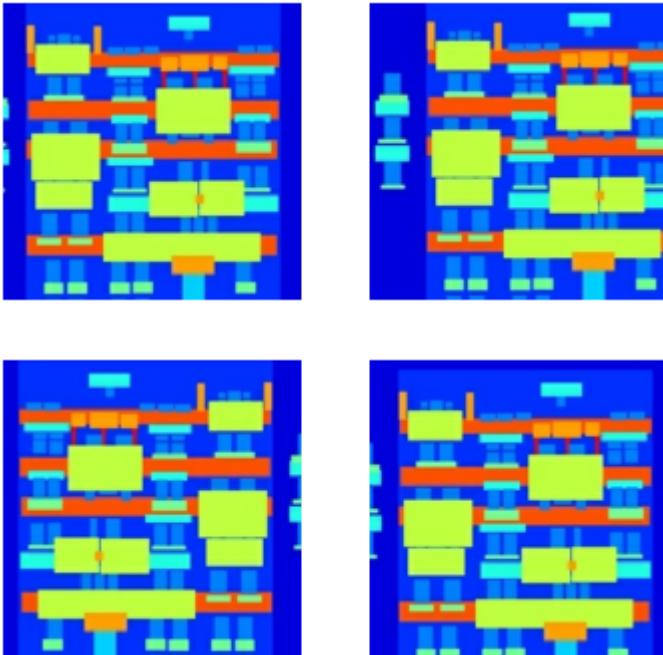
    # Random cropping back to 256x256
    input_image, real_image = random_crop(input_image, real_image, 256, 256)

    if tf.random.uniform(() > 0.5:
```

```
# Random mirroring
input_image = tf.image.flip_left_right(input_image)
real_image = tf.image.flip_left_right(real_image)

return input_image, real_image

plt.figure(figsize=(6, 6))
for i in range(4):
    rj_inp, rj_re = random_jitter(inp, re)
    plt.subplot(2, 2, i + 1)
    plt.imshow(rj_inp / 255.0)
    plt.axis('off')
plt.show()
```



```
def load_image_train(image_file):
    input_image, real_image = load(image_file)
    input_image, real_image = random_jitter(input_image,
    input_image, real_image = normalize(input_image, real

    return input_image, real_image

def load_image_test(image_file):
    input_image, real_image = load(image_file)
    input_image, real_image = resize(input_image, real_im
                                IMG_HEIGHT, IMG_WIDT
    input_image, real_image = normalize(input_image, real

    return input_image, real_image
```

Criar Pipeline

```
train_dataset = tf.data.Dataset.list_files(str(PATH / '
train_dataset = train_dataset.map(load_image_train,
```

```
        num_parallel_calls=tf  
train_dataset = train_dataset.shuffle(BUFFER_SIZE)  
train_dataset = train_dataset.batch(BATCH_SIZE)  
  
try:  
    test_dataset = tf.data.Dataset.list_files(str(PATH /  
except tf.errors.InvalidArgumentError:  
    test_dataset = tf.data.Dataset.list_files(str(PATH /  
test_dataset = test_dataset.map(load_image_test)  
test_dataset = test_dataset.batch(BATCH_SIZE)
```

Construir Gerador

```
OUTPUT_CHANNELS = 3
```

```
def downsample(filters, size, apply_batchnorm=True):  
    initializer = tf.random_normal_initializer(0., 0.02)  
  
    result = tf.keras.Sequential()  
    result.add(  
        tf.keras.layers.Conv2D(filters, size, strides=2,  
                              kernel_initializer=initial  
  
    if apply_batchnorm:  
        result.add(tf.keras.layers.BatchNormalization())  
  
    result.add(tf.keras.layers.LeakyReLU())  
  
    return result  
  
down_model = downsample(3, 4)  
down_result = down_model(tf.expand_dims(inp, 0))  
print (down_result.shape)  
  
(1, 128, 128, 3)
```

```
def upsample(filters, size, apply_dropout=False):  
    initializer = tf.random_normal_initializer(0., 0.02)  
  
    result = tf.keras.Sequential()  
    result.add(  
        tf.keras.layers.Conv2DTranspose(filters, size, stri  
                                      padding='same',  
                                      kernel_initializer=  
                                      use_bias=False))  
  
    result.add(tf.keras.layers.BatchNormalization())  
  
    if apply_dropout:  
        result.add(tf.keras.layers.Dropout(0.5))
```

```
result.add(tf.keras.layers.ReLU())

return result

up_model = upsample(3, 4)
up_result = up_model(down_result)
print (up_result.shape)

(1, 256, 256, 3)

def Generator():
    inputs = tf.keras.layers.Input(shape=[256, 256, 3])

    down_stack = [
        downsample(64, 4, apply_batchnorm=False), # (batch
        downsample(128, 4), # (batch_size, 64, 64, 128)
        downsample(256, 4), # (batch_size, 32, 32, 256)
        downsample(512, 4), # (batch_size, 16, 16, 512)
        downsample(512, 4), # (batch_size, 8, 8, 512)
        downsample(512, 4), # (batch_size, 4, 4, 512)
        downsample(512, 4), # (batch_size, 2, 2, 512)
        downsample(512, 4), # (batch_size, 1, 1, 512)
    ]

    up_stack = [
        upsample(512, 4, apply_dropout=True), # (batch_siz
        upsample(512, 4, apply_dropout=True), # (batch_siz
        upsample(512, 4, apply_dropout=True), # (batch_siz
        upsample(512, 4), # (batch_size, 16, 16, 1024)
        upsample(256, 4), # (batch_size, 32, 32, 512)
        upsample(128, 4), # (batch_size, 64, 64, 256)
        upsample(64, 4), # (batch_size, 128, 128, 128)
    ]

    initializer = tf.random_normal_initializer(0., 0.02)
    last = tf.keras.layers.Conv2DTranspose(OUTPUT_CHANNEL
                                         strides=2,
                                         padding='same'
                                         kernel_initial
                                         activation='ta

    x = inputs

    # Downsampling through the model
    skips = []
    for down in down_stack:
        x = down(x)
        skips.append(x)

    skips = reversed(skips[:-1])

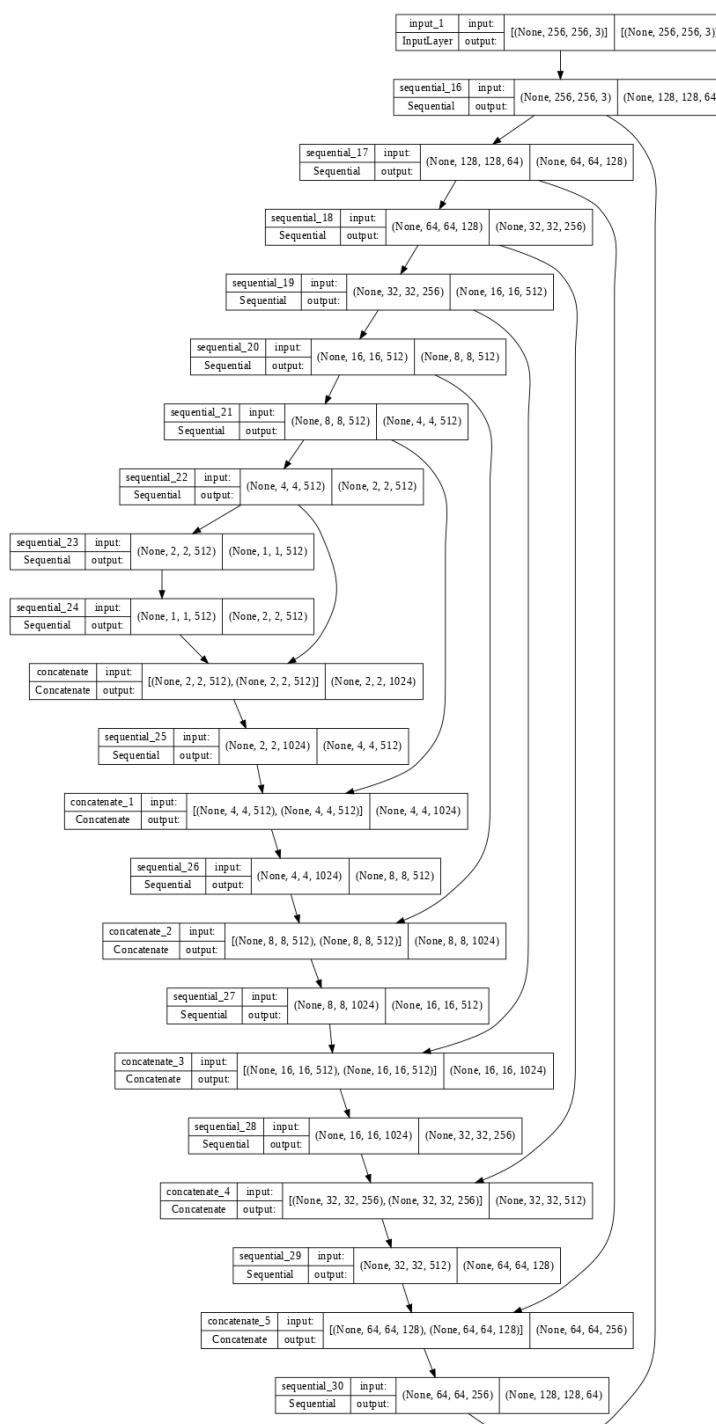
    # Upsampling and establishing the skip connections
    for up, skip in zip(up_stack, skips):
```

```
x = up(x)
x = tf.keras.layers.concatenate([x, skip])

x = last(x)

return tf.keras.Model(inputs=inputs, outputs=x)

generator = Generator()
tf.keras.utils.plot_model(generator, show_shapes=True,
```



```
gen_output = generator(inp[tf.newaxis, ...], training=False)
plt.imshow(gen_output[0, ...])
```

```
WARNING:matplotlib.image:Clipping input data to the  
<matplotlib.image.AxesImage at 0x7f39362a97d0>
```



Definir Perda do Gerador



LAMBDA = 100



```
loss_object = tf.keras.losses.BinaryCrossentropy(from_l
```



```
def generator_loss(disc_generated_output, gen_output, t  
gan_loss = loss_object(tf.ones_like(disc_generated_ou
```

Mean absolute error

```
l1_loss = tf.reduce_mean(tf.abs(target - gen_output))
```

```
total_gen_loss = gan_loss + (LAMBDA * l1_loss)
```

```
return total_gen_loss, gan_loss, l1_loss
```

Construir Discriminador

```
def Discriminator():
```

```
    initializer = tf.random_normal_initializer(0., 0.02)
```

```
    inp = tf.keras.layers.Input(shape=[256, 256, 3], name  
    tar = tf.keras.layers.Input(shape=[256, 256, 3], name
```

```
    x = tf.keras.layers.concatenate([inp, tar]) # (batch
```

```
    down1 = downsample(64, 4, False)(x) # (batch_size, 1  
    down2 = downsample(128, 4)(down1) # (batch_size, 64,  
    down3 = downsample(256, 4)(down2) # (batch_size, 32,
```

```
    zero_pad1 = tf.keras.layers.ZeroPadding2D()(down3) #  
    conv = tf.keras.layers.Conv2D(512, 4, strides=1,  
                                kernel_initializer=init  
                                use_bias=False)(zero_pa
```

```
    batchnorm1 = tf.keras.layers.BatchNormalization()(con
```

```
    leaky_relu = tf.keras.layers.LeakyReLU()(batchnorm1)
```

```
    zero_pad2 = tf.keras.layers.ZeroPadding2D()(leaky_rel
```

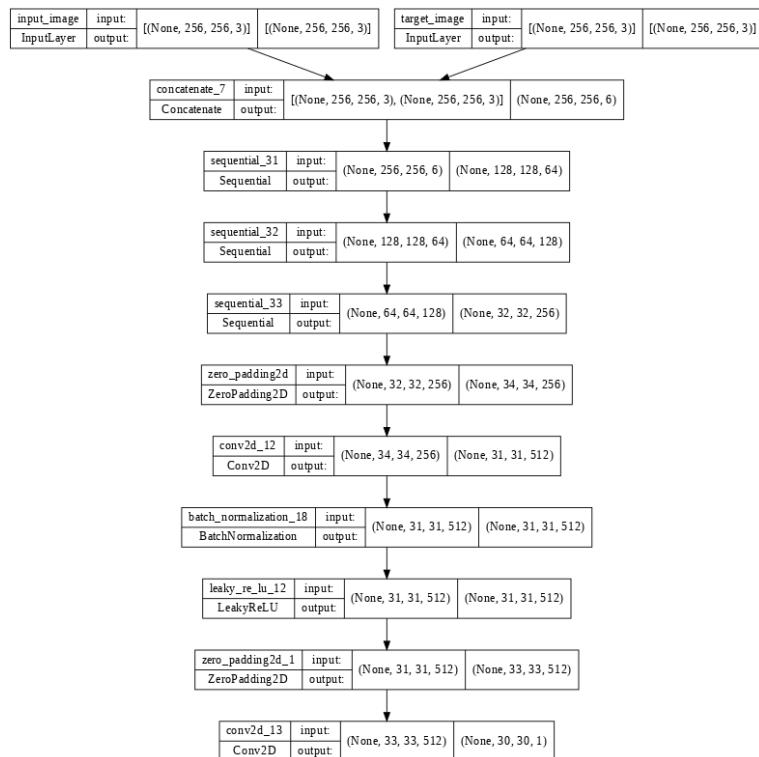
```
    last = tf.keras.layers.Conv2D(1, 4, strides=1,  
                                kernel_initializer=init
```

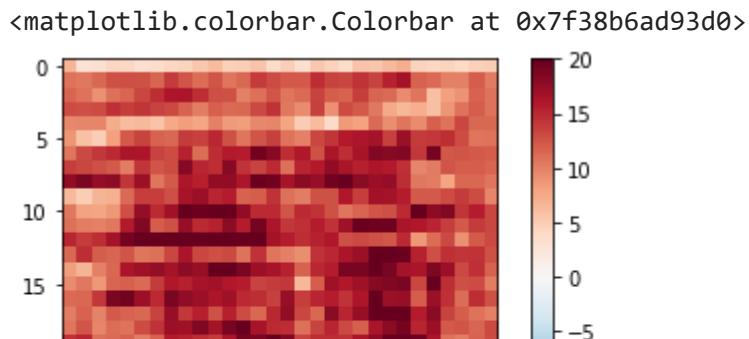
```
    return tf.keras.Model(inputs=[inp, tar], outputs=last
```

```
discriminator = Discriminator()
```

<https://colab.research.google.com/drive/10Xl9SYI88BvZ3YeDaLqGHH0XI6PJcqrl#scrollTo=WVynVI5INsAx&printMode=true>

```
disc_out = discriminator([inp[tf.newaxis, ...], gen_out])
plt.imshow(disc_out[0, ..., -1], vmin=-20, vmax=20, cmap='magma')
plt.colorbar()
```





Definir a perda do discriminador

```
def discriminator_loss(disc_real_output, disc_generated_output):
    real_loss = loss_object(tf.ones_like(disc_real_output),
                           disc_real_output)
    generated_loss = loss_object(tf.zeros_like(disc_generated_output),
                                 disc_generated_output)
    total_disc_loss = real_loss + generated_loss
    return total_disc_loss
```

Definir otimizadores e um protetor de ponto de verificação

```
generator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)
discriminator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)

checkpoint_dir = './training_checkpoints'
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt")
checkpoint = tf.train.Checkpoint(generator_optimizer=generator_optimizer,
                                 discriminator_optimizer=discriminator_optimizer,
                                 generator=generator,
                                 discriminator=discriminator)
```

Gerar Imagens

```
def generate_images(model, test_input, tar):
    prediction = model(test_input, training=True)
    plt.figure(figsize=(15, 15))

    display_list = [test_input[0], tar[0], prediction[0]]
    title = ['Input Image', 'Ground Truth', 'Predicted Image']

    for i in range(3):
        plt.subplot(1, 3, i+1)
        plt.title(title[i])
        # Getting the pixel values in the [0, 1] range to plot
        plt.imshow(display_list[i] * 0.5 + 0.5)
        plt.axis('off')
    plt.show()
```

```
for example_input, example_target in test_dataset.take(
    generate_images(generator, example_input, example_tar
```



Treinamento

```
log_dir="logs/"

summary_writer = tf.summary.create_file_writer(
    log_dir + "fit/" + datetime.datetime.now().strftime("

@tf.function
def train_step(input_image, target, step):
    with tf.GradientTape() as gen_tape, tf.GradientTape()
        gen_output = generator(input_image, training=True)

        disc_real_output = discriminator([input_image, targ
        disc_generated_output = discriminator([input_image,

            gen_total_loss, gen_gan_loss, gen_l1_loss = generat
            disc_loss = discriminator_loss(disc_real_output, di

generator_gradients = gen_tape.gradient(gen_total_los
            generator.trainable_variables)
discriminator_gradients = disc_tape.gradient(disc_los
            discriminator.trainable_variables)

generator_optimizer.apply_gradients(zip(generator_gra
            generator.trainable_variables))
discriminator_optimizer.apply_gradients(zip(discrimin
            discriminator.trainable_variables))

with summary_writer.as_default():
    tf.summary.scalar('gen_total_loss', gen_total_loss,
    tf.summary.scalar('gen_gan_loss', gen_gan_loss, ste
    tf.summary.scalar('gen_l1_loss', gen_l1_loss, step=
    tf.summary.scalar('disc_loss', disc_loss, step=step
```

```
def fit(train_ds, test_ds, steps):
    example_input, example_target = next(iter(test_ds.take(1)))
    start = time.time()

    for step, (input_image, target) in train_ds.repeat():
        if (step) % 1000 == 0:
            display.clear_output(wait=True)

        if step != 0:
            print(f'Time taken for 1000 steps: {time.time() - start:.2f} seconds')

        start = time.time()

        generate_images(generator, example_input, example_target)
        print(f"Step: {step//1000}k")

        train_step(input_image, target, step)

    # Training step
    if (step+1) % 10 == 0:
        print('.', end='', flush=True)

    # Save (checkpoint) the model every 5k steps
    if (step + 1) % 5000 == 0:
        checkpoint.save(file_prefix=checkpoint_prefix)

%load_ext tensorboard
%tensorboard --logdir {log_dir}
```

TensorBoard INACTIVE

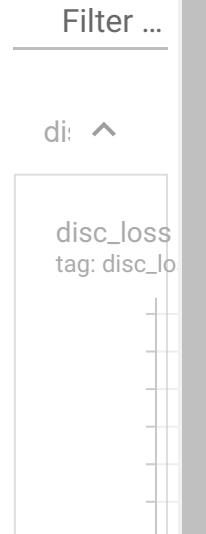
- Show data download links
- Ignore outliers in chart scaling

Tooltip sorting method: default ▾

Smoothing

0,6

Horizontal Axis



```
fit(train_dataset, test_dataset, steps=2000)
```

Time taken for 1000 steps: 98.10 sec



```
display.IFrame(  
    src="https://tensorboard.dev/experiment/1Z0C6FONROa",  
    width="100%",  
    height="1000px")
```

TensorBoard.dev

SEND FEEDBACK

Add a name and description to the experiment to provide more context and details for these results. [Learn more](#)

Created on ...

- Show data download links
- Ignore outliers in chart scaling

Tooltip sorting method: default ▾

Smoothing

 0,6

Horizontal Axis

STEP RELATIVE

WALL

Runs

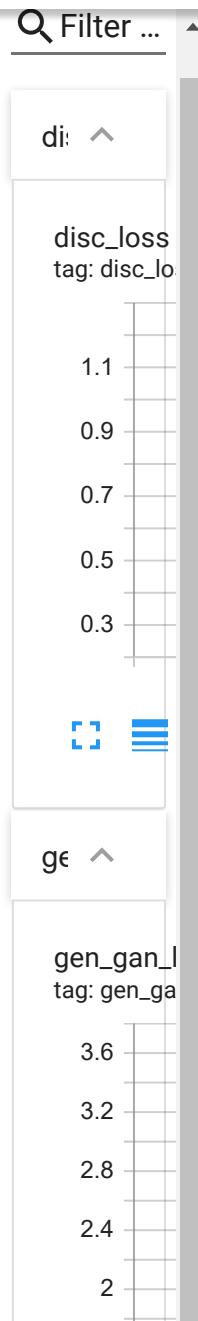
Write a regex to filter runs



TOGGLE ALL RUNS

experiment

IZ0C6FONROaUMfjYkVjJqw



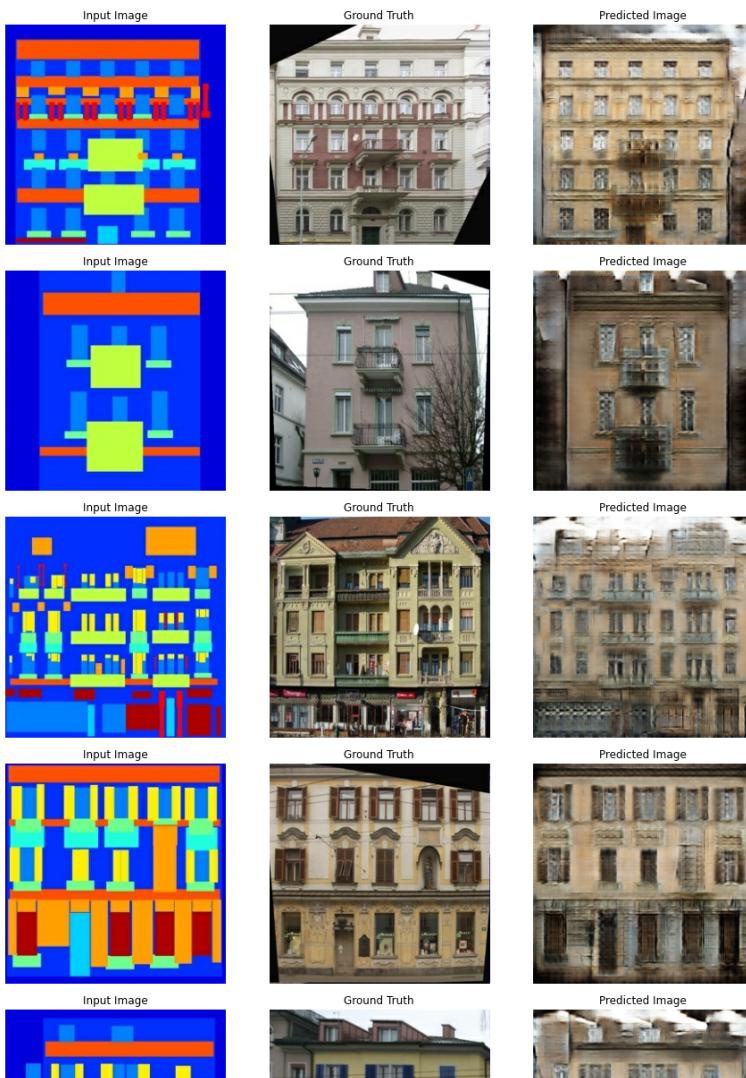
```
ckpt-2.index          ckpt-5.index  
ckpt-3.data-00000-of-00001
```

```
# Restoring the latest checkpoint in checkpoint_dir  
checkpoint.restore(tf.train.latest_checkpoint(checkpoint_dir))
```

```
<tensorflow.python.training.tracking.util.Checkpoint>  
at 0x7f38b6b134d0>
```

Gerar algumas imagens usando o conjunto de teste

```
# Run the trained model on a few examples from the test set  
for inp, tar in test_dataset.take(5):  
    generate_images(generator, inp, tar)
```



Implementar



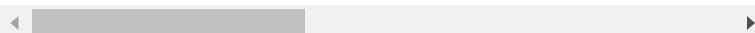
```
# TODO : Criar 3 máscaras e gerar 3 saídas com a pix2pi
```

Baseado nos exemplos do tensorflow [tutorials](#)

```
# Importações
import numpy as np
import os
import cv2
from google.colab.patches import cv2_imshow
```

```
# Ligar ao Drive
from google.colab import drive
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt
```



```
#Ler Imagens (Criei pasta prédios no drive)
root_dir = '/content/drive/My Drive/predios/'
```

```

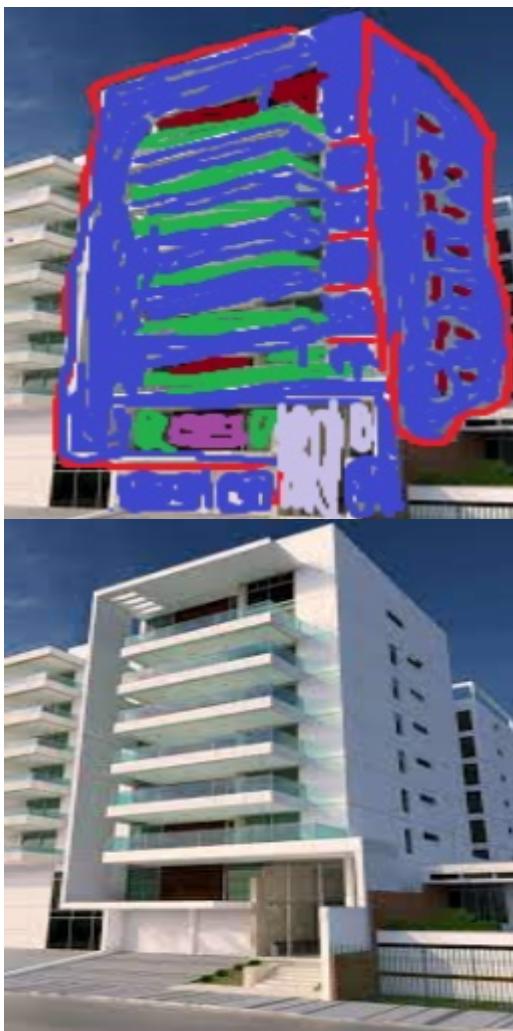
imagens = []
for i in range(1,4):
    input_image = cv2.imread("/content/drive/My Drive/pre
    real_image = cv2.imread("/content/drive/My Drive/pred

    input_image = cv2.resize(input_image, dsize=(256,256)
    real_image = cv2.resize(real_image, dsize=(256,256))

    input_image = np.asarray(input_image)*1.0
    real_image = np.asarray(real_image)*1.0
    """ print(input_image)
    cv2.imshow(input_image)
    cv2.imshow(real_image) """
#input_image, real_image = random_jitter(input_image/
#cv2.imshow( np.asarray(input_image))
#input_image, real_image = normalize(input_image, rea
#cv2.imshow( np.asarray(input_image))
input_image = np.expand_dims(input_image, axis=0)
real_image = np.expand_dims(real_image, axis=0)
imagens["Predio"+str(i)] = [input_image/255, real_im

```

cv2.imshow(imagens["Predio1"])[0][0])
cv2.imshow(imagens["Predio1"])[1][0])



```
# Run the trained model on a few examples from the test
for predio in imagens.keys():
```

```
generate_images(generator, imagens[predio][0], imagens
```

